



Program Code: J620-002-4:2020

Program Name: FRONT-END SOFTWARE DEVELOPMENT

Title : Exe31 - MNIST Handwriting Exercise

Name: Chong Mun Chen

IC Number: 960327-07-5097

Date : 1/8/2023

Introduction : Practising with Neural Network learning model on the MNIST handwritten digit dataset using Tensorflow's Sequential model.

Conclusion : Succeeded in achieving a highly accurate model that predicted the right number from the image.

The Problem: MNIST digit classification

We're going to tackle a classic machine learning problem: MNIST handwritten digit classification. It's simple: given an image, classify it as a digit.



```
In [1]: ▶ import tensorflow as tf
        print(tf.__version__)
```

2.13.0

```
In [2]: ▶ #import all the required libraries
import numpy as np
import mnist
import keras

# The first time you run this might be a bit slow, since the
# mnist package has to download and cache the data.
train_images = mnist.train_images()
train_labels = mnist.train_labels()
test_images = mnist.test_images()
test_labels = mnist.test_labels()
```

Q: What's the dimension of the images data?

```
In [3]: ▶ train_images.shape
```

Out[3]: (60000, 28, 28)

```
In [4]: ▶ test_images.shape
```

Out[4]: (10000, 28, 28)

Q: What's the dimension of the label data?

```
In [5]: ▶ train_labels.size
```

Out[5]: 60000


```
In [9]: train_images.shape
```

```
Out[9]: (60000, 784)
```

```
In [10]: test_images.shape
```

```
Out[10]: (10000, 784)
```

3. Building the Model

Every Keras model is either built using the `Sequential` class, which represents a linear stack of layers, or the functional `Model` class, which is more customizable. We'll be using the simpler `Sequential` model, since our network is indeed a linear stack of layers.

Step: Start by instantiating a `Sequential` model.

- The first two layers have 64 nodes each and use the ReLU activation function.
- The last layer is a Softmax output layer with 10 nodes, one for each class.

Q: what's the correct input shape for your input layer?

```
In [11]: from keras.models import Sequential
         from keras.layers import Dense

         # Define the model
         model = Sequential([
             Dense(64, activation='relu'),
             Dense(64, activation='relu'),
             Dense(10, activation='softmax'),
         ])
```

4. Compiling the Model

Before we can begin training, we need to configure the training process. We decide 3 key factors during the compilation step:

- The optimizer. We'll stick with a pretty good default: the Adam gradient-based optimizer. Keras has many other optimizers you can look into as well.
- The loss function. Since we're using a Softmax output layer, we'll use the Cross-Entropy loss. Keras distinguishes between `binary_crossentropy` (2 classes) and `categorical_crossentropy` (>2 classes), so we'll use the latter
- A list of metrics. Since this is a classification problem, we'll just have Keras report on the accuracy metric.

Step: Compile the model using the above options - `adam`, `categorical_crossentropy`, `accuracy` as metrics

```
In [12]: model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)
```

5. Training the Model

Training a model in Keras literally consists only of calling `fit()` and specifying some parameters. There are a lot of possible parameters, but we'll only manually supply a few:

- The training data (images and labels), commonly known as X and Y, respectively.
- The number of epochs (iterations over the entire dataset) to train for.
- The batch size (number of samples per gradient update) to use when training.

Step: set epochs to a suitable number, and `batch_size = 32`

```
In [13]: from keras.models import Sequential
    from keras.layers import Dense
    from keras.utils import to_categorical

    # Train the model.
    model.fit(
        train_images,
        to_categorical(train_labels),
        epochs=5,
        batch_size=32,
    )
```

```
Epoch 1/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.3581 -
accuracy: 0.8921
Epoch 2/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.1923 -
accuracy: 0.9419
Epoch 3/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.1525 -
accuracy: 0.9522
Epoch 4/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.1278 -
accuracy: 0.9594
Epoch 5/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.1106 -
accuracy: 0.9649
```

```
Out[13]: <keras.src.callbacks.History at 0x1a20624de70>
```

Q: Do you run into any problem? Why?

```
In [ ]: 
```

Q: what's your achieved accuracy?

6. Testing the Model

Step: Evaluating the model by testing against the test data

```
In [14]: ▶ model.evaluate(
            test_images,
            to_categorical(test_labels)
        )

313/313 [=====] - 1s 1ms/step - loss: 0.1349 - a
ccuracy: 0.9598

Out[14]: [0.1348690688610077, 0.9598000049591064]
```

7. Using the Model

Now that we have a working, trained model, let's put it to use. The first thing we'll do is save it to disk so we can load it back up anytime.

Step: save the model using the `save_weights` function

```
In [15]: ▶ model.save_weights('model.h5')

In [16]: ▶ from tensorflow.keras.models import Sequential
            from tensorflow.keras.layers import Dense

            # Build the model.
            model = Sequential([
                Dense(64, activation='relu', input_shape=(784,)),
                Dense(64, activation='relu'),
                Dense(10, activation='softmax'),
            ])

            # Load the model's saved weights.
            model.load_weights('model.h5')
```

8. Predict

Using the trained model to make predictions is easy: we pass an array of inputs to `predict()` and it returns an array of outputs. Keep in mind that the output of our network is 10 probabilities (because of softmax), so we'll use `np.argmax()` to turn those into actual digits.

```
In [17]: ▶ # Predict on the first 5 test images.
predictions = model.predict(test_images[:5])

# Print our model's predictions.
print(np.argmax(predictions, axis=1)) # [7, 2, 1, 0, 4]

# Check our predictions against the ground truths.
print(test_labels[:5]) # [7, 2, 1, 0, 4]

1/1 [=====] - 0s 197ms/step
[7 2 1 0 4]
[7 2 1 0 4]
```

Note: What's the difference between `model.save_weights` and `model.save`? -

[https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save\(\)%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else](https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save()%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else)
([https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save\(\)%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else](https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save()%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else)).

This exercise is adapted from <https://victorzhou.com/blog/keras-neural-network-tutorial/>
(<https://victorzhou.com/blog/keras-neural-network-tutorial/>).

Challenge 1:

Retrain your model by using different network depths - what will you conclude?

```
In [18]: ▶ model = Sequential([
    Dense(64, activation = 'relu', input_shape = (784,)),
    Dense(64, activation = 'relu'),
    Dense(64, activation = 'relu'),
    Dense(64, activation = 'relu'),
    Dense(10, activation = 'softmax')
])
```

Challenge 2:

Retrain your model by using different activation (other than ReLU) - what differences does it make?

```
In [19]: ▶ model = Sequential([
    Dense(64, activation = 'sigmoid', input_shape = (784,)),
    Dense(64, activation = 'sigmoid'),
    Dense(10, activation = 'softmax'),
])
```

Challenge 3:

```
In [20]: ▶ from tensorflow.keras.optimizers import Adam

model.compile(
    optimizer = Adam(lr = 0.005),
    loss = 'categorical_crossentropy',
    metrics = ['accuracy']
)

model.fit(
    train_images,
    to_categorical(train_labels),
    epochs = 5,
    batch_size = 32,
    validation_data = (test_images, to_categorical(test_labels))
)
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.Adam`.

```
Epoch 1/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.5713 -
accuracy: 0.8494 - val_loss: 0.2594 - val_accuracy: 0.9247
Epoch 2/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.2316 -
accuracy: 0.9315 - val_loss: 0.1908 - val_accuracy: 0.9434
Epoch 3/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.1694 -
accuracy: 0.9492 - val_loss: 0.1792 - val_accuracy: 0.9443
Epoch 4/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.1345 -
accuracy: 0.9601 - val_loss: 0.1415 - val_accuracy: 0.9581
Epoch 5/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.1133 -
accuracy: 0.9661 - val_loss: 0.1169 - val_accuracy: 0.9631
```

Out[20]: <keras.src.callbacks.History at 0x1a206a210c0>

Challenge 4:

How will you load your saved weights to use it in a separate code? Upload your saved model/weights, and compare your model/weights with a model/weights from one of your classmate's.

```
In [21]: ▶ model.save_weights('model.h5')
```

Challenge 5:

How can you load any image from the data set and let your model (or your classmate's) to predict the image?


```
In [22]: ► predictions = model.predict(test_images[:5])

print(np.argmax(predictions, axis = 1))
print(test_labels[:5])

1/1 [=====] - 0s 48ms/step
[7 2 1 0 4]
[7 2 1 0 4]
```