



Big Data Analytic Platform

Instructor, Nero Chan Zhen Yu



What is Apache Hadoop



- A software framework for storing, processing, and analyzing “big data”
 - Distributed
 - Scalable
 - Fault--tolerant
 - Open source

Something about Apache Hadoop

- Open source software
- Around 60 committers from companies and volunteer developers
 - Cloudera, Yahoo!, Facebook, Apple and more
- A large software ecosystem



Why do we need something like Hadoop?

- To solve problems that exist in traditional large-scale analytics systems
 - Computation has been processor (and memory) –bound
 - Can't just keep buying bigger computers
 - Programming for traditional distributed systems is complex
 - Synchronization for data exchange? Partial failure problem
 - Distributed systems
 - Programming complexity
- Solutions before Hadoop - ?
 - New approach needed!

Hence we need a system that can....

- Support partial failure
 - Failure of a component will not cause complete failure of the entire system
- Data recoverability
 - Automatic restart task and recover lost data, or better – no lost of data even when there is a component failure
- Components can “come and go”
 - No restart of system is required
- Consistency
 - Outcome of the job is not affected even if there is a component failure
- Scalability
 - Additional load will not bring down the system
 - Increase resources can be added flexibly



**Forward
School**

A little history

- Hadoop is based on Google File System in 2003, later MapReduce in 2004
- Core concept:
 - distribute data across the nodes in a distributed system
 - processing is only done on local machine/node itself
 - Programming – high-level code
 - Communication between nodes – minimal
 - Replication of data – redundancy to improve availability and reliability

Hadoop Core Components



MapReduce



Hadoop Distributed File System

Hadoop
Core
Components


Hadoop Distributed File System





- Based on a presented white paper on Google File System (GFS)
- Initially developed as a storage infrastructure for Apache Nutch web search engine project
- Some characteristics:
 - Extremely fault-tolerant
 - Can hold large number of datasets
 - Provides redundant storage for massive amounts of data

How files are stored in HDFS


- Files are split into blocks of 64MB or 128MB (Typically)
- Data is distributed across many machines at load time
 - The blocks will be stored on at least 3 machines across the cluster
 - Provides local processing, especially for efficient MapReduce processing



Distribution of data



Name Node to manage metadata



File access


- To users, they are just access a file
 - Via FsShell command (hadoop fs)
 - Java API
 - Ecosystem projects (Flume, Sqoop and Hue)
- Hadoop/HDFS manages the access to the respective blocks

Hadoop MapReduce

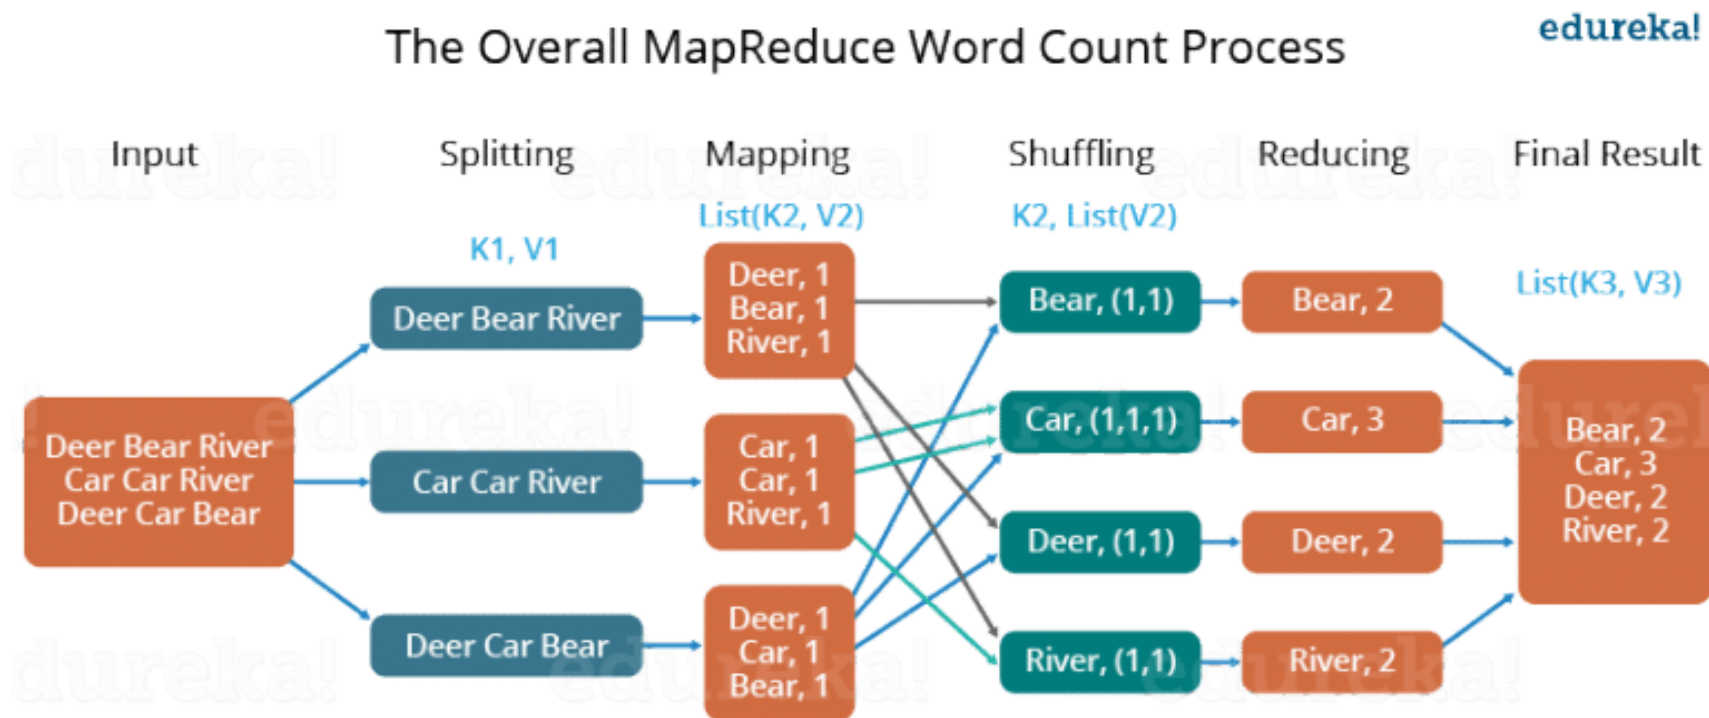
- Hadoop's implementation of MapReduce
- MapReduce
 - A method for distributing a task across multiple nodes in a cluster
 - Each node processes data (remember the blocks?) stored at that particular node
 - Components - Mapper, Reducer and Shuffle and Sort
 - Terminology
 - A job – a full programme
 - A task – execution of a task (a single Mapper/Reducer) on a block
 - A task attempt – a particular instance of an attempt to execute a task

Hadoop MapReduce

- Mapper
 - Process data on a single HDFS block at the node where the block is stored
 - Usually the first part of the manipulation of <key,value> in a certain processing needs
- Shuffle and sort
 - Sorts and consolidates intermediate data from all mappers
- Reducer
 - Process data from Shuffle and Sort to produce the final output (second part of the processing/manipulation of data)



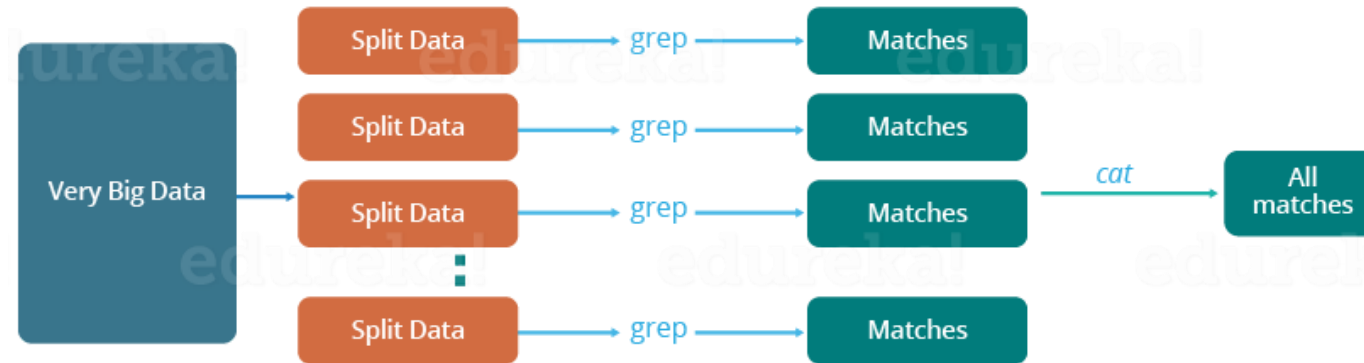
A word count example



Comparison

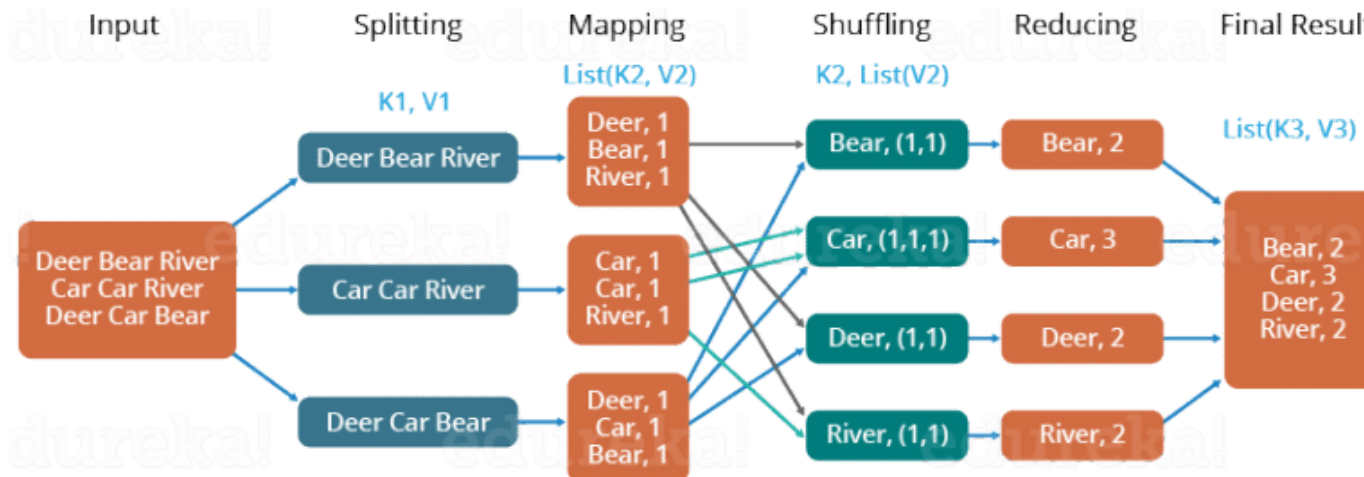
The Traditional Way

edureka!



The Overall MapReduce Word Count Process

edureka!



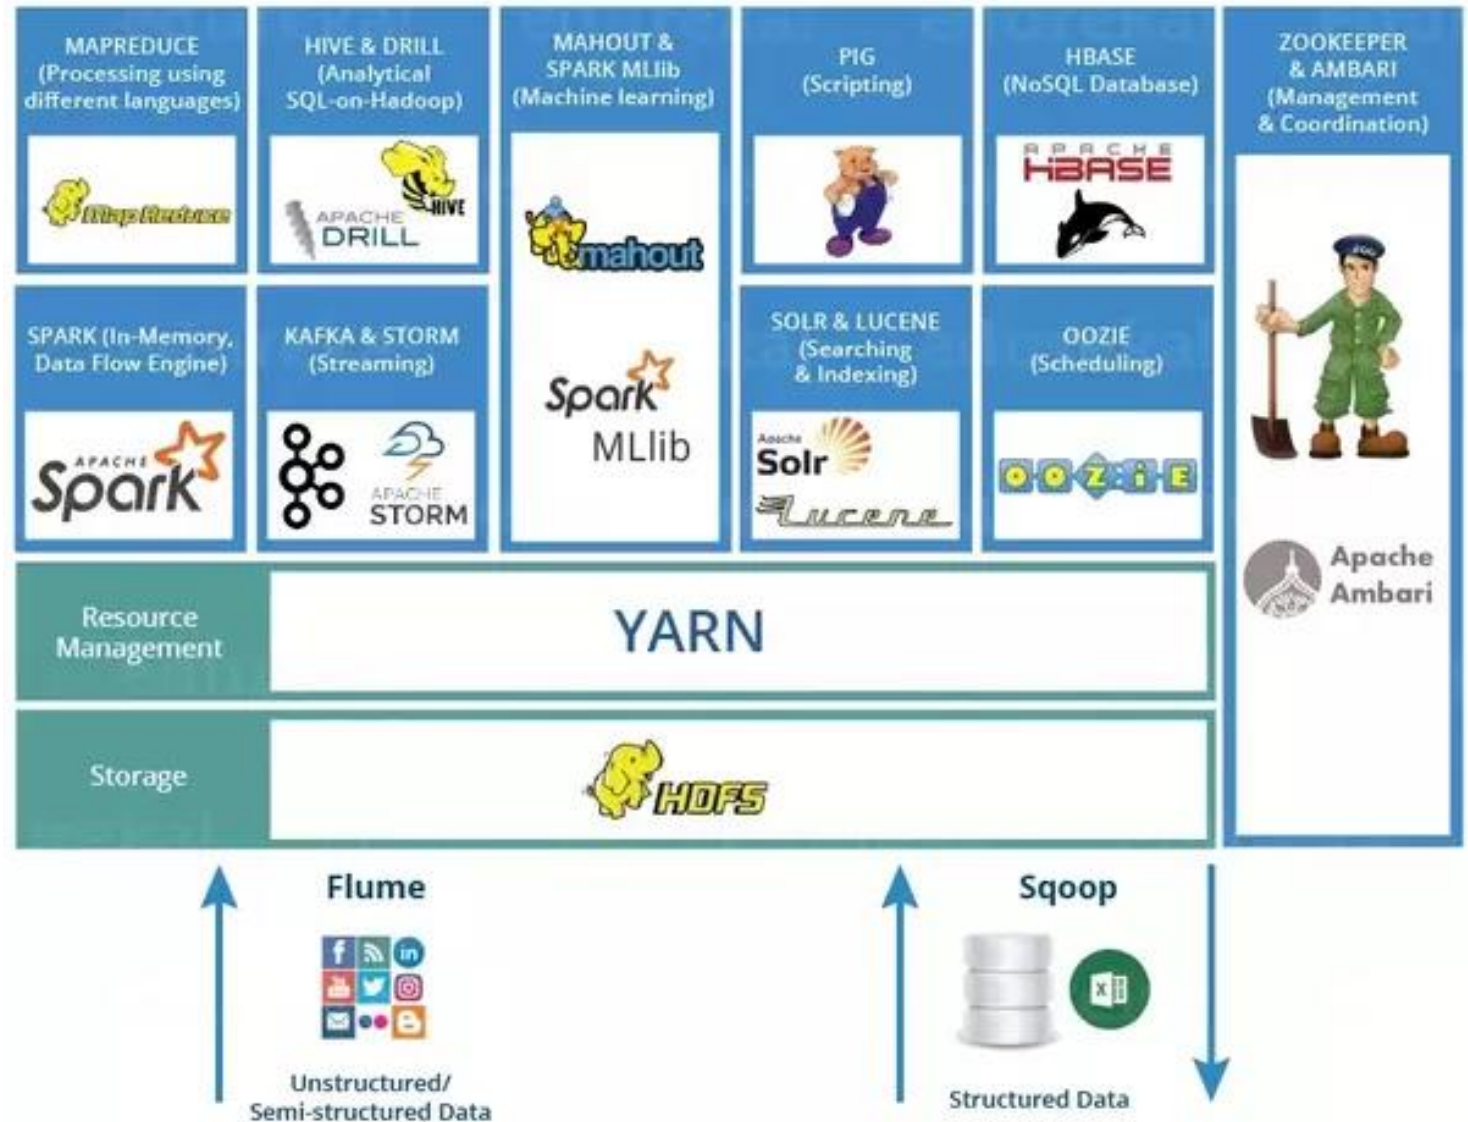
If word count is boring....

- Processing web log
- Find Top Ten records
- Find distinct values

Hadoop Ecosystem

Ecosystem....

- Built on HDFS
- Built on HDFS and MapReduce
- Designed to integrate with or support Hadoop



Data Storage: HBase



- Hadoop Database
- NoSQL Datastore
- Can store more than Petabytes of data (Massive amount of data)
- High write throughput rate
- Handles sparse data well – no wasted spaces for empty columns in a row
- Limitations
 - Only optimized for row look up by key – no FULL queries like SQL
 - No transaction: single row operations only
 - Only the key is indexed

Comparison with Traditional RDBMS

	RDBMS	Hbase
Data layout	Row-oriented	Column-oriented
Transactions	Yes	Single row only
Query language	SQL	Get/put/scan
Security	Authentication/Authorization	Kerberos
Indexes	Any column	Only row-key
Max data size	TBs	PB+
Read/write throughput (queries per second)	Thousands	Millions

When you should use Hbase

- Use HDFS if
 - You only append to your dataset (no random write)
 - You usually read the whole dataset (no random read)
- Use HBase if
 - You need random write and/or read
 - You do thousands of operations per second on TB+ of data
- Use an RDBMS if
 - Your data fits on one big node
 - You need full transaction support
 - You need real-time query capabilities


Data Integration: Flume



- What is Flume?
 - A service to move large amounts of data in real time
 - Example: storing log files in HDFS
- Flume imports data into HDFS as it is generated
 - Instead of batch-processing it later
 - For example, log files from a Web server
- Flume is
 - Distributed
 - Reliable and available
 - Horizontally scalable
 - Extensible

Flume – High level overview

- Source may be files, logs, stdout or custom
- Scalable throughput to write in parallel
- Store in any format
 - Text, compressed, binary or custom




Data Integration: Sqoop



- Exchanging data with RDBMS
- Sqoop transfer data between RDBMS and HDFS very efficiently
- Supports JDBC, ODBC, and other specific databases
- Custom connectors
 - MySQL, Postgres, Netezza, Teradata, Oracle
- Not open source but free to use

Data Processing: Spark

- a unified analytics engine for large-scale data processing
- Spark demonstrates high performance for batch and streaming data
- Uses DAG scheduler
- Supports Java, Scala, Python, R and SQL
- Runs on Hadoop, and more – Apache Mesos, Kubernetes, standalone



Spark

- Originally developed in UC Berkely's AMPLab
- Benefits over MapReduce
 - Speed – way faster than MapReduce
 - Better suited for iterative algorithms
 - Can hold intermediate data in RAM, resulting in much better performance
 - Easier API
 - Supports real-time streaming data processing

Data Analysis: Hive, Pig and Impala

- MapReduce is powerful, but hard to code/master
- High level programming to perform MapReduce
 - Hive and Pig – Languages for querying and manipulating data
 - Support/leverage on existing skillsets
 - SQL
 - Programmers
- Open source Apache projects
- Interpreter turns queries into MapReduce jobs

Hive



- HiveQL – An SQL-like interface to Hadoop/MapReduce

```
SELECT * FROM purchases WHERE price > 10000 ORDER BY  
storeid
```

Pig



Apache Pig

- A scripting dataflow language (called Pig Latin) for transforming large data sets

```
purchases = LOAD "/user/dave/purchases" AS (itemID,  
                                             price, storeID, purchaserID);  
bigticket = FILTER purchases BY price > 10000;  
...
```

Comparison

	Hive	Pig
Language	HiveQL (SQL-like)	Pig Latin (dataflow language)
Schema	Table definitions stored in a metastore	Schema optionally defined at runtime
Programmatic access	JDBC, ODBC	PigServer (Java API)

Impala

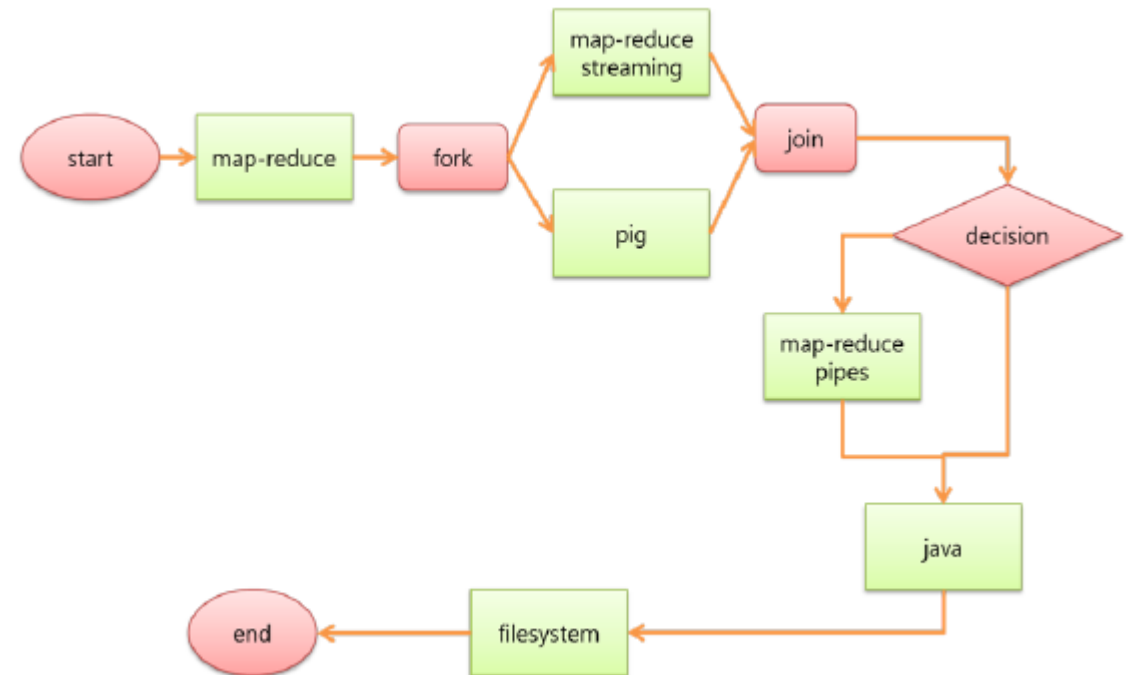
- High performance SQL engine for vast amounts of data
 - Similar query language to HiveQL
 - 10 to 50+ times faster than Hive, Pig or MapReduce
- Impala runs on Hadoop clusters
 - Data stores in HDFS
 - Does not use MapReduce
- 100% opensource but developed by Cloudera



Workflow Engine: Oozie



- Workflow engine for MapReduce jobs
- Defines dependencies between jobs
- Ensure jobs are submitted in the correct sequences





Machine Learning: Mahout

- Mahout is a Machine Learning library written in Java
- Use for
 - Collaborative filtering
 - Clustering
 - Classification
- Why use Hadoop for Machine Learning?
 - Because of the ability to store large amount of data – and the machine learning outcome may have the advantage